

Implementation Blueprint: Instant Domain AI Visibility Scorecard and llms.txt Validator (AIQualityHQ.com)

1. Market Opportunity Analysis

The global digital discovery ecosystem is undergoing a structural paradigm shift from traditional keyword-indexed retrieval to generative, agentic answer synthesis. Between early 2025 and 2026, user query volume routed through artificial intelligence engines—such as OpenAI ChatGPT Search, Perplexity, Google AI Overviews, Anthropic Claude, and Gemini—surged by 42.8%, expanding from 15.6 billion to over 27.4 billion monthly search interactions. This transition has destabilized conventional Search Engine Optimization (SEO), diminishing organic click-through rates by 30% to 60% as users receive direct, synthesized answers in place of traditional blue links.

The evolution of answer engine technology has created Generative Engine Optimization (GEO) and Answer Engine Optimization (AEO) as core business functions. While the overall market for search and visibility tools stands at \$68.7 billion globally, the specific market for AI Search Analytics and GEO software attracted over \$300 million in venture capital funding between mid-2025 and spring 2026.

Market Segment	Valuation / Volume	Growth Profile & Key Dynamics
Total Addressable Market (TAM)	\$68.7 Billion	Global SEO, digital marketing, and web analytics market undergoing AI disruption.
Serviceable Addressable Market (SAM)	\$3.2 Billion	Dedicated Generative Engine Optimization (GEO) & AI Answer Monitoring software.
Serviceable Obtainable Market (SOM)	\$450 Million	Developer, SMB, and agency self-serve AI audit tooling and scorecards.

Market growth is accompanied by significant structural inefficiencies across existing tooling solutions:

- **Excessive Price Barriers in Enterprise Incumbents:** Market leaders such as Profound—valued at \$1 billion after raising \$155 million—gate key analytics, agent traffic logs, and multi-engine tracking behind enterprise contracts or high monthly tiers ranging from \$99/month for single-engine tracking to \$499+/month. Mid-market teams, smaller agencies, and indie developers are effectively priced out of continuous AI visibility monitoring.
- **Friction-Heavy Onboarding Workflows:** Existing visibility suites require manual prompt setup, complex workspace configurations, and multi-hour data processing windows before delivering baseline insights. The market lacks an instant, unauthenticated single-input audit engine capable of returning diagnostic visibility scores in under three seconds.
- **Disconnected Technical & Context Standards:** While the llms.txt standard proposed by Jeremy Howard in September 2024 has gained adoption across developer-focused

ecosystems like Anthropic, Vercel, Stripe, and Mintlify, existing platforms evaluate llms.txt in isolation. They fail to correlate machine-readable Markdown files with broader entity authority, JSON-LD schema parsing, and crawler access permissions.

AIQualityHQ.com directly captures this market opportunity by delivering a frictionless domain audit engine that merges llms.txt validation, entity authority scoring, and multi-engine AI visibility tracking into a single unified platform.

Strategic Recommendations

- Position AIQualityHQ.com as the universal "Speedtest + Google PageSpeed Insights for AI Search," offering zero-friction public audits to capture top-of-funnel demand.
- Leverage an unauthenticated, instant scorecard model to build an authoritative domain index that fuels programmatic SEO traffic growth.

2. Customer Research and Segment Analysis

The customer base for AI search optimization spans technical developers, growth marketers, and enterprise strategy teams. Each segment exhibits distinct buying triggers, technical capabilities, and budgetary constraints.

Customer Segment	Core Pain Points	Buying Triggers	Search Intent & Query Examples	Monthly Budget	Existing Tool Stack	Key Complaints with Existing Tools
SaaS Founders	Zero visibility into how LLMs describe their software; losing organic referral traffic to competitors.	Competitor cited as "Best SaaS for X" in ChatGPT while own domain is omitted.	"how to show up in ChatGPT", "llms.txt generator", "GEO tools".	\$49 - \$149	Google Analytics, Search Console, Ahrefs	High enterprise pricing, overly complex setup, lack of clear technical recommendations.
SEO & Marketing Agencies	Clients demanding AI visibility reports; inability to audit or bill for GEO services at scale.	Losing clients due to declining traditional search traffic; need white-label AI audits.	"white label AI visibility tool", "agency GEO report", "AEO audit tool".	\$199 - \$499	Semrush, Ahrefs, BrightEdge, Screaming Frog	Enterprise tools lack white-label PDF/URL exports; per-prompt pricing scales exponentially.
Developers & Tech Leads	Documentation blocked by JS rendering;	Deploying API docs; implementing llms.txt	"llms.txt specification", "validate llms.txt",	\$0 - \$49 (or API)	Vercel, GitHub, Mintlify, Cloudflare	No programmatic validation API; manual

Customer Segment	Core Pain Points	Buying Triggers	Search Intent & Query Examples	Monthly Budget	Existing Tool Stack	Key Complaints with Existing Tools
	tokens wasted on HTML boilerplate when LLMs read docs.	standard; building internal RAG bots.	"markdown docs for LLM".		Workers	Markdown syntax checking is error-prone.
AI Startups & Developer Platforms	AI agents fetching outdated product context; hallucinated pricing and feature specs.	Product launch on Product Hunt; AI agents failing to parse product API.	"fix ChatGPT hallucination about brand", "optimize site for AI agents".	\$99 - \$299	Firecrawl, OpenAI API, Anthropic SDK	Tools track mentions but do not diagnose <i>why</i> crawlers misread documentation.
E-commerce Brands	Omitted from ChatGPT Shopping and Perplexity product recommendations.	Direct drop in top-of-funnel organic product discovery sales.	"ChatGPT shopping visibility", "Perplexity e-commerce SEO".	\$149 - \$399	Shopify, Klaviyo, Google Merchant Center	Existing tools focus strictly on B2B documentation rather than product schema extraction.
Publishers & Media	AI crawlers scraping content without attribution or referral traffic.	Traffic drop due to Google AI Overviews; blocking training bots vs search bots.	"block GPTBot allow OAI-SearchBot", "robots.txt for AI crawlers".	\$99 - \$299	WordPress, Ghost, Cloudflare	Complex crawler rule management; lack of log analysis for AI user-agents.

Strategic Recommendations

- Build specific landing page entry points tailored to developers (/llms-txt/validator), agencies (/solutions/agencies), and SaaS founders (/score/domain).
- Ensure the free tier resolves immediate developer technical pain points (llms.txt syntax validation) while gating agency reporting features behind paid tiers.

3. Competitor Research and Strategic Positioning

The competitive environment divides into enterprise-focused monitoring platforms, mid-market analytics tools, and traditional SEO suites attempting to bolt on generative search tracking modules.

| Platform | Entry Pricing | AI Engine Coverage | Key Strengths | Core Weaknesses | Missing Features | AIQualityHQ Strategy | | :--- | :--- | :--- | :--- | :--- | :--- | :--- | | **Profound** | \$99/mo (Lite), \$399/mo (Growth), Enterprise Custom | ChatGPT, Perplexity, Gemini, Claude, Copilot, Grok, Meta, DeepSeek | Deep conversation analytics, agent traffic logs, massive runway (\$155M raised). | High cost, complex setup, engine coverage gated by tier. | Unauthenticated instant scorecard, native llms.[span_15](start_span)[span_15](end_span)[span_35](start_span)[span_35](end_span)txt validator. | Undercut entry price with an instant, zero-friction public audit flywheel. | | **Peec AI** | ~€89 - €120/mo | ChatGPT, Perplexity, Gemini, Copilot, Google AIO | Clean UX, prompt suggestions, competitive sentiment benchmarking. | Lacks deep technical/crawl audit layer, prompt volume limits. | Technical schema auditing, llms.txt parser. | Combine technical readiness audits with citation tracking. | | **Otterly AI** | \$25 - \$29/mo | ChatGPT, Gemini, Copilot, Perplexity, Google AIO | Affordable entry point, simple keyword-to-prompt converter. | Surface-level analytics, limited prompt frequency, basic reporting. | Deep agent crawl log analysis, entity graph mapping. | Provide comprehensive technical readiness evaluation and code fixes. | | **ZipTie** | ~\$58.65/mo | Google AI Overviews, ChatGPT, Perplexity | Deep visual answer-level auditing for AI Overviews. | Limited model coverage, no conversation-level tracking. | Bypasses llms.txt and machine-readability optimizations. | Merge Google AIO analysis with LLM Markdown optimization. | | **Semrush / Ahrefs AI** | Bundled into \$139+/mo base plans | Basic ChatGPT / Brand Mention tracking | Mass distribution across existing traditional SEO user base. | Shallow feature depth, delayed data updates, high user seat fees. | Purpose-built GEO architecture centered on token efficiency. | Focus on specialized GEO technical audits that legacy suites ignore. |


Strategic Recommendations

- Differentiate AIQualityHQ.com as the single platform combining technical crawlability auditing (robots.txt AI user-agent checking), llms.txt validation, and real-time multi-engine citation scoring.
- Maintain a transparent self-serve pricing model (\$49 Pro, \$199 Agency) to capture mid-market users frustrated by enterprise sales calls.

4. Product Vision and User Journey

AIQualityHQ.com operates a product-led growth (PLG) flywheel driven by an unauthenticated, zero-friction audit experience.

User Journey Progression:



```

[3. Interactive Scorecard Output (Score: 0-100)]
  └─ Diagnostic
Sub-Sc[span_174] (start_span) [span_174] (end_span) [span_178] (start_span)
[span_178] (end_span) [span_182] (start_span) [span_182] (end_span)ores
(Readiness, GEO, Citation, Technical)
  └─ Interactive llms.txt Syntax & Token Analysis
  └─ High-Impact Vulnerability Warnings
      |
      ▼
[4. Conversion Trigger & Paid Workflow]
  └─ Claim Domain & Generate One-Click Fixes
  └─ Upgrade to
P[span_103] (start_span) [span_103] (end_span) [span_123] (start_span) [span
_123](end_span)ro/Agency for Continuous Monitoring

```

End-to-End User Experience Lifecycle

1. **Discovery & Top-of-Funnel Entry:** A user lands on AIQualityHQ.com via a programmatic SEO scorecard page (e.g., /score/stripe.com) or directly via the homepage. The hero UI features a clean, single-input search interface: *"Enter Domain to Audit AI Visibility & llms.txt Compliance"*.
2. **Instant Asynchronous Execution:** Upon domain submission, Cloudflare Workers initiate parallel edge requests:
 - *Probe A:* Fetches /llms.txt, /llms-full.txt, and checks .md endpoint availability.
 - *Probe B:* Parses /robots.txt for explicit rules governing AI search bots (OAI-SearchBot, Claude-SearchBot, PerplexityBot) versus training bots (GPTBot, ClaudeBot, Google-Extended).
 - *Probe C:* Scrapes root HTML to extract JSON-LD entities (Organization, SoftwareApplication, [span_139](start_span)[span_139](end_span)[span_150](start_span)[span_150](end_span)Product) and semantic structure.
 - *Probe D:* Queries cached SERP/LLM endpoints to measure immediate citation frequency across category prompts.
3. **Interactive Scorecard Presentation:** Within three seconds, the user receives an overall AI Visibility Scorecard (0–100 Rating) broken down into sub-scores. Critical vulnerabilities—such as blocking OAI-SearchBot or publishing an invalid llms.txt file—are flagged with immediate remediation steps.
4. **Actionable Remediation & Tooling:** The platform offers direct code fixes, including buttons to *"Generate Optimized llms.txt"* or *"Copy Corrected robots.txt Rules"*.
5. **Account Creation & Monetization:** To track historical score trends, set up automated weekly alerts, compare share of voice against competitors, or export white-label PDF reports, users transition into paid tiers.

Strategic Recommendations

- Ensure the initial scorecard load completes in under 3 seconds to maximize conversion and user delight.

- Provide one-click downloadable fixes (llms.txt, robots.txt, JSON-LD snippets) directly on the free scorecard to showcase immediate utility.

5. Feature Roadmap and Prioritization Matrix

Features are structured to prioritize instant diagnostic utility for early acquisition while building defensible reporting moats over time.

Feature Name	Category	Purpose	User Benefit	Technical Complexity	Business Value	Revenue Potential
Instant Scorecard Engine	MVP	Computes an instant 0–100 AI Visibility Score for any domain.	Instant diagnostic rating without requiring account registration.	Low (Edge Worker + KV Cache)	Critical (Primary viral growth hook)	High (Top-of-funnel driver)
llms.txt Syntax Validator	MVP	Validates Markdown structure against Jeremy Howard's specification.	Identifies broken links, syntax errors, and missing blockquotes.	Low (Regex + AST Parser)	High (Establishes technical authority)	Medium (Drives developer acquisition)
AI Crawler Permission Auditor	MVP	Parses robots.txt for AI search versus training bot access rules.	Prevents unintended exclusion from ChatGPT and Perplexity search indexes.	Low (Robots.txt parser)	High (Immediate technical utility)	Low (Core free feature)
Automated llms.txt Builder	MVP	Scrapes site content and generates a fully compliant llms.txt file.	One-click generation of machine-readable documentation.	Medium (DOM extraction + LLM summary)	High (High lead-capture conversion)	Medium (Included in Free/Pro)
Multi-Engine Citation Tracker	Growth	Tracks brand mentions across ChatGPT, Perplexity, Gemini, Claude, and Google AIO.	Quantifies brand citation share across target AI engines.	Medium (DataForSEO SERP / LLM APIs)	Critical (Core recurring subscription value)	High (Primary driver for Pro/Agency)
Competitor Share of	Growth	Benchmarks target	Reveals competitor	Medium (Parallel	High (Essential	High (Drives upgrades to

Feature Name	Category	Purpose	User Benefit	Technical Complexity	Business Value	Revenue Potential
Voice (SoV)		domain against named competitors across prompts.	ranking advantages in AI search answers.	prompt queries)	feature for agencies)	Agency tier)
Schema & Entity Authority Audit	Growth	Audits JSON-LD entities and Knowledge Graph connectivity.	Enhances machine-readability for RAG pipelines.	Medium (JSON-LD parser + Wikidata API)	High (Deep technical differentiation)	Medium (Drives Pro upgrades)
Token Waste & Context Analyzer	Premium	Evaluates HTML-to-token bloat and Markdown parsing efficiency.	Minimizes token expenditure for AI agents fetching docs.	High (Tokenizer + AST transformer)	Medium (Appeals to developers)	Medium (Gated Pro feature)
White-Label PDF & URL Reports	Premium	Generates branded client audit reports for marketing agencies.	Allows agencies to resell GEO diagnostic services.	Medium (Puppeteer / Edge PDF rendering)	High (Unlocks agency sales)	High (Drives Agency \$199/mo tier)
Real-Time AI Bot Log Monitor	Enterprise	Tracks server logs for direct OAI-SearchBot and ClaudeBot visits.	Verifies live agent crawling activity on specific paths.	High (Cloudflare Logpush / Ingestion pipeline)	High (Defensible enterprise moat)	Very High (Enterprise \$499+/mo tier)


Strategic Recommendations

- Focus initial sprint cycles strictly on MVP features to launch a functional public scorecard within 30 days.
- Gate competitive Share of Voice benchmarking and white-label reporting behind the Agency tier to maximize ARPU.

6. AI Visibility Scoring Framework

The overall **AI Visibility Score** (S_{AIV}) is a weighted composite metric ranging from 0 to 100, calculated from ten specialized sub-scores:

$$S_{\text{AIV}} = \sum_{i=1}^{10} w_i \cdot S_i = w_1 S_{\text{GEO}} + w_2 S_{\text{AIR}} + w_3 S_{\text{CIT}} + w_4 S_{\text{ENT}} + w_5 S_{\text{LLM}} + w_6 S_{\text{SCH}} + w_7 S_{\text{TRU}} + w_8 S_{\text{CNT}} + w_9 S_{\text{CRW}} + w_{10} S_{\text{RNK}}$$



Comprehensive Mathematical Formulation of Sub-Scores

1. AI Crawlability Score (S_{CRW} , Weight: $w_1 = 0.12$)

Evaluates /robots.txt directives to confirm whether search indexing agents (OAI-SearchBot, Claude-SearchBot, PerplexityBot) and user agents are permitted access.

$$S_{\text{CRW}} = 100 \times \left(0.60 \times \frac{N_{\text{allow}}}{N_{\text{total_search}}} + 0.20 \times \frac{N_{\text{allowed_user}}}{N_{\text{total_user}}} + 0.20 \times \mathbb{I}(\text{WAF Block} = \text{False}) \right)$$

2. llms.txt Score (S_{LLM} , Weight: $w_2 = 0.10$)

Measures compliance with Jeremy Howard's specification, validating header structure, blockquote summaries, link health, and companion /llms-full.txt files.

$$S_{\text{LLM}} = 20 \times \mathbb{I}(\text{Exists}) + 20 \times \mathbb{I}(\text{Valid H1/Blockquote}) + 20 \times \left(\frac{L_{\text{valid}}}{L_{\text{total}}} \right) + 20 \times \mathbb{I}(\text{llms-full.txt Exists}) + 20 \times \mathbb{I}(\text{Zero Syntax Errors})$$

3. AI Readiness Score (S_{AIR} , Weight: $w_3 = 0.08$)

Evaluates server latency, JavaScript dependencies, and clean Markdown availability via .md URL extensions.

$$S_{\text{AIR}} = 40 \times \mathbb{I}(\text{No JS Required}) + 30 \times \max\left(0, 1 - \frac{\text{TTFB}}{1000\text{ms}}\right) + 30 \times \mathbb{I}(\text{.md Endpoint Available})$$

4. Entity Authority Score (S_{ENT}), Weight: $w_4 = 0.10$)

Verifies brand presence across foundational knowledge databases, including Wikidata, Google Knowledge Graph, Wikipedia, and Crunchbase.

$$S_{\text{ENT}} = 35 \cdot \mathbb{I}(\text{Wikidata ID}) + 35 \cdot \mathbb{I}(\text{Google Knowledge Graph ID}) + 30 \cdot \min\left(1, \frac{\text{Verified SameAs Links}}{5}\right)$$

5. Schema Score (S_{SCH}), Weight: $w_5 = 0.10$)

Audits JSON-LD structured data implementation (Organization, SoftwareApplication, Product, FAQPage, TechArticle).

$$S_{\text{SCH}} = 50 \cdot \mathbb{I}(\text{Valid Core JSON-LD}) + 30 \cdot \mathbb{I}(\text{Zero Syntax Errors}) + 20 \cdot \mathbb{I}(\text{Nested Entity Context})$$

6. Trust Score (S_{TRU}), Weight: $w_6 = 0.10$)

Evaluates domain security, SSL certificates, privacy/terms documentation, author schemas, and trusted seed domain backlinks.

$$S_{\text{TRU}} = 30 \cdot \mathbb{I}(\text{HTTPS/SSL}) + 30 \cdot \mathbb{I}(\text{E-E-A-T Schema Present}) + 40 \cdot \min\left(1, \frac{\text{Domain Rank}}{100}\right)$$

7. Citation Score (S_{CIT}), Weight: $w_7 = 0.15$)

Measures the percentage of targeted industry prompts where the domain is cited as an explicit source across ChatGPT, Perplexity, Gemini, and Claude.

$$S_{\text{CIT}} = 100 \cdot \left(\frac{\sum_{m \in \text{Models}} \text{Citations}_m}{N_{\text{prompts}} \cdot N_{\text{models}}} \right)$$

8. GEO Score (S_{GEO}), Weight: $w_8 = 0.15$)

Calculates relative Share of Voice (SoV) and brand mention density compared to top three direct industry competitors across AI answers.

$$S_{\text{GEO}} = 100 \cdot \left(\frac{\text{Brand Mentions}}{\text{Brand Mentions} + \sum \text{Competitor Mentions}} \right)$$

9. Content Score (S_{CNT}), Weight: $w_9 = 0.05$)

Evaluates informational density, answer directness (FAQ structure), and lack of marketing jargon.

$$S_{\text{CNT}} = 50 \cdot (\text{Flesch-Kincaid Clarity Factor}) + 50 \cdot \left(\frac{\text{Facts and Statistics Count}}{\text{Total Word Count}} \right)$$

10. Ranking Position Score (S_{RNK}), Weight: $w_{10} = 0.05$)

Measures ordinal list placement when cited in recommendation lists (e.g., Position #1 vs

Position #5).

$$S_{\text{RNK}} = 100 \times \left(1 - \frac{\bar{P}}{1 + N_{\text{ma}}(\text{start_span})(\text{end_span})(\text{span_256})(\text{start_span})(\text{span_256})(\text{end_span})(\text{span_259})(\text{start_span})(\text{span_259})(\text{end_span}) \times \text{list_items}}} \right)$$

Score Anti-Manipulation & Fraud Detection Safeguards

To prevent artificial score inflation via keyword stuffing or cloaking:

- **Synthetic Citation Verification:** Cross-validates generated citations against live SERP responses to eliminate LLM hallucination false positives.
- **Edge Cloaking Detection Engine:** Compares edge HTTP responses against generic browser user-agents and AI crawler user-agents to detect differential content serving.
- **Uninformative Keyword Penalty:** Applies negative scaling factors to domains utilizing repetitive, uninformative keyword blocks to game semantic matching.

Strategic Recommendations

- Display the sub-score breakdown clearly on scorecards so users understand *why* their domain received a specific rating.
- Highlight critical quick wins (e.g., updating robots.txt or creating llms.txt) that deliver immediate score gains.

7. llms.txt Research and Validation Engine

The llms.txt standard, introduced by Jeremy Howard (Answer.AI) in September 2024, provides structured Markdown summaries to guide language models through web content without consuming excessive context window capacity.

Structure of /llms.txt File:

```
|— # Project Name (Required H1)
|— > Short Summary Blockquote (Recommended 1-2 Sentences)
|— System Context & Operating Guidance (Optional Text)
|— ## Core Section Header (H2 Group)
|   |— [Title] (https://domain.com/page): Markdown link + context
note
|   |— [Title] (https://domain.com/page.md): Direct Markdown target
link
|— ## Optional Section Header (H2 Group - Skippable Content)
    |— [Secondary Link] (URL): Context for deep agent retrieval
```

Technical Validation Logic & Specification Rules

The validation engine enforces strict parsing rules based on the official specification at llmstxt.org:

1. **Location & Path Standards:** Must be served at /llms.txt (or a specific subpath for sub-projects) with a
Con[span_112](start_span)[span_112](end_span)[span_132](start_span)[span_132](end_

span)tent-Type of text/markdown or text/plain; charset=utf-8.

2. **Structural Syntax Breakdown:**

- **Required Single H1:** Exactly one # Title element defining the brand or project name.
- **Required Blockquote:** A blockquote > directly beneath the H1 providing a concise summary.
- **Structured H2 File Lists:** Section groupings (## Documentation, ## API Refere[span_391](start_span)[span_391](end_span)[span_397](start_span)[span_397](end_span)[span_403](start_span)[span_403](end_span)n) containing bulleted Markdown hyperlinks.
- **Optional Section Convention:** An ## Optional header designating secondary links that AI agents can bypass when context windows are constrained.

3. **Hyperlink Syntax Rules:** * Links must adhere strictly to Markdown formatting:

[Name](URL).

- Explanatory notes must follow a colon delimiter: [Name](URL): Detailed not[span_441](start_span)[span_441](end_span)[span_444](start_span)[span_444](end_span)[span_447](start_span)[span_447](end_span)es.
- Links pointing directly to clean .md endpoints are prioritized over HTML pages.

Diagnostic Error Matrix and Auto-Remediation Code Fixes

Diagnostic Error Code	Detection Logic	Severity	Automated Remediation Recommendation
ERR_MISSING_FILE	HTTP 404/500 at /lms.txt.	Critical	Run Auto-Builder to generate and publish an initial /lms.txt file.
ERR_INVALID_H1	Missing # header or multiple H1 elements present.	High	Restructure Markdown header to feature exactly one # Brand Name line.
ERR_NO_BLOCKQUOTE	Absence of > blockquote directly beneath H1.	Medium	Insert a concise 150-character domain summary in a > blockquote.
ERR_BROKEN_LINK	Hyperlink within file returns 4xx/5xx HTTP status code.	Critical	Remove or update broken URL references.
ERR_TOKEN_WASTE	Links point to heavy HTML pages instead of clean .md endpoints.	Medium	Append .md extensions or serve pure Markdown versions of linked pages.
ERR_MISSING_FULL	HTTP 404 at /lms-full.txt when /lms.txt exists.	Low	Concatenate primary documentation pages into a single /lms-full.txt file.
[span_312](start_span)	Narrative text exceeds	Medium	Trim system context

Diagnostic Error Code	Detection Logic	Severity	Automated Remediation Recommendation
span_312](end_span)[span_321](start_span)[span_321](end_span)ERR_VERBOSE_CONTEXT	500 tokens before file lists.		prose; move non-essential guidance to linked files.

Strategic Recommendations

- Build the validator as an open-source npm package (@aiqualityhq/llmstxt-validator) to gain developer adoption and natural backlinks.
- Provide a one-click button on audit reports that automatically fixes syntax errors and outputs downloadable Markdown.

8. Comprehensive AI Visibility Audit System

The diagnostic engine executes a multi-layered audit across technical, structural, and entity signals.

Multi-Layered AI Visibility Audit Matrix:

1. AI CRAWLER ACCESS LAYER
 - └─ robots.txt AI Search Bot Permissions (OAI-SearchBot, Claude-SearchBot)
 - └─ robots.txt AI Training Bot Permissions (GPTBot, ClaudeBot, Google-Extended)
 - └─ Firewall & WAF Edge Inspection (Verifies [span_243](start_span) [span_243](end_span) [span_245](start_span) [span_245](end_span)non-blocking status)
2. MACHINE-READABLE ARCHITECTURE LAYER
 - └─
/ll[span_203](start_span) [span_203](end_span) [span_210](start_span) [span_210](end_span)ms.txt &
/llms-fu[span_383](start_span) [span_383](end_span) [span_385](start_span) [span_385](end_span) [span_387](start_span) [span_387](end_span)ll.txt Spec
Complia[span_338](start_span) [span_338](end_span) [span_345](start_span) [span_345](end_span)nce Verification
 - └─ Markdown Mirror Endpoint Availability (.md URL Fallbacks)
 - └─ JavaScript Rendering Dependency Evaluation
3. ENTITY & STRUCTURED DATA LAYER
 - └─ JSON-LD Schema (Organization, TechA[span_229](start_span) [span_229](end_span) [span_236](start_span) [span_236](end_span)rticle, Product, FAQPage)
 - └─ Knowledge Graph Connectivity (Wikidata ID, Google [span_230](start_span) [span_230](end_span) [span_237](start_span) [span_237](end_span)

237](end_span)Knowledge Graph)
└─ E-E-A-T & Author Identity Verification Markup

4. GENERATIVE CITATION LAYER
- └─ Multi-Model Citation Frequency (ChatGPT, Perplexity, Gemini, Claude)
 - └─ Competitor Share of Voice (SoV) & Ordinal List Position
 - └─ Brand Mention Sentiment Analysis

Proprietary Audit Signals Competitors Overlook

1. **AI User-Agent Disconnect Identification:** Competitors treat all AI bots identically. The diagnostic engine differentiates between **Training Crawlers** (GPTBot, ClaudeBot) and **Search/Answer Crawlers** (OAI-S[span_140](start_span)[span_140](end_span)[span_151](start_span)[span_151](end_span)earchBot, Claude-SearchBot, PerplexityBot). Blocking training bots while allowing search bots protects intellectual property without sacrificing generative search visibility.
2. **Token-to-Information Efficiency Ratio:** Quantifies the token cost of stripping HTML/CSS/JS boilerplate relative to useful text context. Clean, token-efficient sites receive higher priority in RAG retrieval pipelines.
3. **Semantic Markdown Heading Depth:** Verifies that content utilizes structured sequential headings (#, ##, ###) rather than visually styled <div> elements, ensuring parseability by LLM AST chunkers.

Strategic Recommendations

- Alert users immediately if their robots.txt configuration accidentally blocks search indexing bots while intending to block training bots.
- Provide an actionable breakdown showing how much token context users can save by serving clean Markdown versions of key documentation.

9. UI/UX Architecture and Interface Specifications

The design framework prioritizes speed, high data density, and clear diagnostic recommendations.

1. Global Homepage & Instant Audit Search Interface

UI Component	Structural Layout & Design Elements
Hero Section	Clean centered layout with headline: <i>"Measure & Maximize Your Domain's AI Visibility"</i> . Large single-input search bar accepting domain entries (e.g., stripe.com) with a prominent CTA button: <i>"Run Instant Audit"</i> .
Trust Badges	Horizontal ticker showing compliance checks:

UI Component	Structural Layout & Design Elements
	[✓] Verified against llms.txt standard [✓] Multi-Engine AI Crawler Check [✓] JSON-LD Schema Diagnostic.
Live Audit Feed	Real-time horizontal ticker displaying recent public audits: github.com (94/100) linear.app (88/100) vercel.com (96/100).

2. Interactive Audit Report Dashboard

UI Component	Structural Layout & Design Elements
Header Card	Displays Target Domain, Overall Score Gauge (e.g., 88/100 - Grade: A), Last Checked timestamp, and CTAs: [Export White-Label PDF] [Claim Domain].
Sub-Score Grid	4-column diagnostic cards: AI Crawlability (100/100) llms.txt Score (80/100) Schema & Entity (90/100) Citation Share (82/100).
Remediation Panel	Highlighted action alerts: [!] [span_205](start_span)[span_205](end_span)[span_212](start_span)[span_212](end_span)Warning: /llms-f[span_422](start_span)[span_422](end_span)[span_426](start_span)[span_426](end_span)ull.txt missing. Click to generate. [!] Action: 2 links in /llms.txt return 404 redirects..
Citation Matrix	Multi-engine citation performance table detailing brand mention position across ChatGPT, Perplexity, Claude, and Gemini.

3. Pricing Page Interface

UI Component	Structural Layout & Design Elements
Tier Grid	4-column pricing layout displaying Free (\$0), Pro (\$49/mo), Agency (\$199/mo), and Enterprise (\$499+/mo) with annual billing toggles offering a 20% discount.
Feature Comparison	Matrix listing monitored prompts, domain audits, white-label exports, team seats, and server log monitoring features.

Strategic Recommendations

- Ensure error states (e.g., invalid domain, server timeout) display helpful fallback guidance rather than blank screens.
- Design empty states on claimed user dashboards with explicit step-by-step onboarding prompts to encourage initial setup.

10. Programmatic SEO Engine

The programmatic SEO (pSEO) strategy drives organic user acquisition by indexing scorecards, comparison reports, and validator tools for thousands of domains.

Programmatic SEO Information Architecture:

- └─ /score/[domain] (e.g., /score/stripe.com)
- └─ /compare/[domA]-vs-[domB] (e.g., /compare/stripe.com-vs-paypal.com)
- └─ /industry/[category] (e.g., /industry/developer-tools)
- └─ /llms-txt/[domain] (e.g., /llms-txt/stripe.com)

Dynamic Schema & Indexing Infrastructure

Every programmatic page dynamically renders schema markup to maximize traditional and generative search indexing:

```
{
  "@context": "https://schema.org",
  "@type": "SoftwareApplication",
  "name": "AI Visibility Scorecard for stripe.com",
  "operatingSystem": "Web",
  "applicationCategory": "SEO/GEO Diagnostic Tool",
  "offers": {
    "@type": "Offer",
    "price": "0.00",
    "priceCurrency": "USD"
  },
  "aggregateRating": {
    "@type": "AggregateRating",

[span_114] (start_span) [span_114] (end_span) [span_134] (start_span) [span_
134] (end_span) "ratingValue": "88",
    "bestRating": "100",
    "worstRating": "0",
    "ratingCount": "1"
  }
}
```

Projected Organic Traffic & Backlink Growth

Programmatic SEO Asset	Target Page Count (Year 1)	Estimated Monthly Organic Traffic	Target Backlink Generation
Domain Scorecards (/score/)	50,000 pages	250,000 visits/mo	1,500 domain backlinks
Versus Comparisons	20,000 pages	180,000 visits/mo	800 domain backlinks

Programmatic SEO Asset	Target Page Count (Year 1)	Estimated Monthly Organic Traffic	Target Backlink Generation
(/compare/)			
llms.txt Validators (/llms-txt/)	50,000 pages	120,000 visits/mo	3,000 developer backlinks
Industry Benchmarks (/industry/)	500 pages	50,000 visits/mo	600 press backlinks
Total pSEO Network	120,500 pages	600,000 visits/mo	5,900 high-authority links

Strategic Recommendations

- Implement dynamic HTML caching at the Cloudflare Edge to serve pSEO pages in under 100 milliseconds.
- Automatically generate canonical links pointing to primary domain scorecards to avoid duplicate content penalties.

11. Generative Engine Optimization (GEO) Strategy

To ensure AIQualityHQ.com itself becomes the authoritative answer when users query ChatGPT, Perplexity, Gemini, Claude, or Google AI Overviews regarding AI visibility or llms.txt validation, the domain executes a targeted GEO framework.

AIQualityHQ Self-Optimization Framework:

1. STRUCTURED NATIVE ENDPOINTS
 - └ Host native /llms.txt & /llms-full.txt files
 - └ Serve JSON-LD SoftwareApplication & Dataset markup
2. MULTI-PLATFORM ENTITY SEEDING
 - └ Maintain active profiles on G2, GitHub, Product Hunt, Reddit
 - └ Secure verified entity links in Wikidata & Knowledge Graph
3. EXTRACTABLE CONTENT CHUNKS
 - └ Use direct Q&A headers (## How to validate llms.txt?)
 - └ Publish quarterly data reports cited by tech publications

Execution Tactics

1. **Serve Flawless Native Context Endpoints:** AIQualityHQ.com hosts exemplary native /llms.txt and /llms-full.txt files, serving as an industry reference model for AI search crawlers.
2. **Entity Seeding Across High-Citation Sources:** Generative engines rely heavily on trusted platforms for knowledge retrieval. The platform maintains verified profiles, documentation, and data feeds across GitHub, Product Hunt, Reddit (r/SEO, r/SaaS), Crunchbase, and G2.
3. **Direct Q&A Content Formatting:** Content across AIQualityHQ.com uses concise, direct semantic HTML (<h2>Question</h2><p>Direct 1-sentence answer.</p>), providing easily extractable text chunks for Retrieval-Augmented Generation (RAG) pipelines.

Strategic Recommendations

- Monitor citations of AIQualityHQ.com weekly using internal prompt tracking to ensure accurate representation across all LLM engines.
- Publish open benchmark datasets on GitHub to earn citations in AI research papers and developer guides.

12. Production-Ready Database Design

The platform uses Supabase (PostgreSQL) optimized for high-throughput edge reads and secure tenant isolation.

```
-- Core Domain ENUM Types
CRE[span_488] (start_span) [span_488] (end_span)ATE TYPE user_role AS
ENUM ('user', 'admin', 'agency_member');
CREATE TYPE sub_plan AS ENUM ('free', 'pro', 'agency', 'enterprise');
CREATE TYPE audit_status AS ENUM ('pending', 'processing',
'completed', 'failed');

-- Users & Profiles
CREATE TABLE users (
    id UUID PRIMARY KEY
DEFAUL[span_261] (start_span) [span_261] (end_span) [span_263] (start_span)
[span_263] (end_span)T gen_random_uuid(),
    email TEXT UNIQUE NOT NULL,
    created_at TIMESTAMPTZ DEFAULT NOW() NOT NULL,
    updated_at TIMESTAMPTZ DEFAULT NOW() NOT NULL
);

-- Organizations (Supports Agency Multi-Tenancy)
CREATE TABLE organizations (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    name TEXT NOT NULL,
    owner_id UUID REFERENCES users(id) ON DELETE CASCADE NOT NULL,
    plan sub_plan DEFAULT 'free' NOT NULL,
    stripe_customer_id TEXT,
    stripe_subscription_id TEXT,
    created_at TIMESTAMPTZ DEFAULT NOW() NOT NULL
);

-- Monitored Domains
CREATE TABLE domains (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    org_id UUID REFERENCES organizations(id) ON DELETE CASCADE,
    domain_name TEXT UNIQUE NOT NULL,
    is_public BOOLEAN DEFAULT true NOT NULL,
    created_at TIMESTAMPTZ
```

```

DEFAULT[span_79](start_span)[span_79](end_span)[span_82](start_span)[s
pan_82](end_span)[span_85](start_span)[span_85](end_span) NOW() NOT
NULL
);

```

-- AI Visibility Score Reports

```

CREATE TABLE ai_visibility_reports (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    domain_id UUID REFERENCES domains(id) ON DELETE CASCADE NOT NULL,
    overall_score INT NOT NULL,
    crawlability_score INT NOT NULL,
    llms_txt_score INT NOT NULL,
    readiness_score INT NOT NULL,
    entity_score INT NOT NULL,
    schema_score INT NOT NULL,
    trust_score INT NOT NULL,

```

```

citation_sc[span_115](start_span)[span_115](end_span)[span_135](start_
span)[span_135](end_span)ore INT NOT NULL,
    geo_score INT NOT NULL,
    status audit_status DEFAULT 'pending' NOT NULL,
    raw_payload JSONB,
    created_at TIMESTAMPTZ DEFAULT NOW() NOT NULL
);

```

-- Monitored Prompts

```

CREATE TABLE monitored_prompts (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    domain_id UUID REFERENCES domains(id) ON DELETE CASCADE NOT NULL,
    prompt_text TEXT NOT NULL,
    category TEXT,
    created_at TIMESTAMPTZ DEFAULT NOW() NOT NULL
);

```

-- Citation Tracking History

```

CREATE TABLE citation_mentions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    prompt_id UUID REFERENCES monitored_prompts(id) ON DELETE CASCADE
NOT NULL,
    model_name TEXT NOT NULL,
    is_cited BOOLEAN DEFAULT false NOT NULL,
    ordinal_position INT,
    sentiment_score NUMERIC(3,2),
    snippet_text TEXT,
    checked_at TIMESTAMPTZ DEFAULT NOW() NOT NULL
);

```

-- Speed Optimization Indexes

```
CREATE INDEX idx_domains_name ON domains(domain_name);
CREATE INDEX idx_reports_domain_created ON
ai_visibility_reports(domain_id, created_at DESC);
CREATE INDEX idx_citations_prompt_checked ON
citation_mentions(prompt_id, checked_at DESC);
```


Strategic Recommendations

- Utilize PostgreSQL Row-Level Security (RLS) policies on organizations and domains tables to enforce secure multi-tenant data isolation.
- Apply automated partitioning to citation_mentions tables by month to maintain database query speed as tracking volume grows.

13. Technical Architecture and Edge Stack

The system leverages a serverless edge architecture built on Next.js, Cloudflare, and Supabase.

Technical Infrastructure Stack:



Stack Components and Functionality

- **Frontend & Edge Serving:** Next.js App Router deployed on Cloudflare Pages, offering global edge rendering and near-zero cold starts for programmatic SEO routes.
- **Edge Processing Workers:** Cloudflare Workers parse HTML, validate /lms.txt, and check /robots.txt permissions at edge nodes in milliseconds.
- **Asynchronous Queue Pipeline:** Heavy jobs—such as querying multi-engine SERPs or

executing LLM evaluations—are offloaded to Cloudflare Queues, preventing API timeout errors during real-time user sessions.

- **Global Caching Layer:** Cloudflare KV caches audit scorecards with a 7-day TTL, serving public programmatic SEO requests instantly while minimizing outbound API expenses.

Strategic Recommendations

- Deploy strict rate-limiting policies (e.g., 10 free audits per IP per hour) at the Cloudflare WAF level to prevent scraping abuse.
- Implement health monitoring via Sentry to track edge execution exceptions in real time.

14. API Research and Data Sourcing Integration

To balance latency, financial cost, and data reliability, API dependencies are structured hierarchically.

API Provider	Operational Purpose	Pricing Structure	Rate Limits	Fallback / SLA Alternative
DataForSEO SERP API	Fetches real-time Google AI Overviews and SERP citations.	\$0.0006 - \$0.0012 per query	2,000 requests/min	Serper.dev / ValueSERP
OpenAI API (GPT-4o-mini)	Evaluates brand citation context and extracts sentiment.	\$0.15 / 1M input tokens	10,000 RPM	DeepSeek-V3 / Groq API
Anthropic API (Claude 3.5 Sonnet)	Audits Markdown quality and validates documentation structures.	\$3.00 / 1M input tokens	1,000 RPM	Claude 3.5 Haiku
Perplexity API	Queries real-time Perplexity search citations and source links.	\$0.005 per query	500 RPM	Headless Browser Scraper
Firecrawl API	Scrapes JavaScript-heavy web applications into Markdown.	\$0.002 per crawl	1,000 requests/min	ScrapingBee / Puppeteer Core

Strategic Recommendations


- Route all high-volume sentiment checks through low-cost models (GPT-4o-mini) to keep operational overhead low.
- Maintain secondary API keys and automatic failover logic to prevent service disruptions if a primary API provider experiences downtime.

15. Cost Optimization and AI Economics

Operating a high-margin SaaS platform requires strict separation between deterministic logic (costing \$0 in AI tokens) and probabilistic LLM evaluation logic.

Cost Containment Routing Pipeline:

Incoming Domain Request



Cost Reduction Framework

1. **Deterministic Processing First:** 70% of the overall audit score (Crawlability, llms.txt syntax, Schema validation, Readiness) is calculated using pure JavaScript AST/Regex parsing at the Cloudflare Edge, incurring \$0 in LLM API fees.
2. **Aggressive Edge KV Caching:** All domain scorecards are cached in Cloudflare KV for 7 days. If multiple users query stripe.com, only the initial request incurs external API costs.
3. **Prompt Evaluation Batching:** Instead of making individual API calls per prompt, the engine combines up to 5 target prompts into a single structured JSON system call to GPT-4o-mini, reducing input token overhead by 80%.

Strategic Recommendations

- Require user authentication for deep multi-engine citation checks to prevent unauthorized bot traffic from inflating API costs.
- Set strict daily spending caps on OpenAI and DataForSEO billing dashboards to avoid unexpected cost overruns.

16. Pricing and Revenue Strategy

The pricing structure accommodates self-serve individual users while capturing agency and enterprise upside.

Subscription Tier	Monthly Price	Annual Price	Monitored Prompts	Included Audits	Included Core Features	Primary Target Audience
Free / Tool	\$0	\$0	1 Prompt	3 Audits / mo	Instant Scorecard, llms.txt Validator, Basic robots.txt check.	Developers, casual users.
Pro	\$49 / mo	\$39 / mo	25 Prompts	20 Audits / mo	Automated llms.txt Builder, Full Multi-Engine Citation Tracking, Schema Fixer.	SaaS Founders, Solo Marketers.
Agency	\$199 / mo	\$159 / mo	150 Prompts	Unlimited	White-Label PDF/URL Reports, Share of Voice Competitor Benchmarking, 5 Seats.	SEO & Marketing Agencies.
Enterprise	\$499+ / mo	\$399+ / mo	Custom	Custom	Server Log push integration for AI Bots, Custom API, Dedicated Slack Channel.	Mid-market & Enterprise Brands.

Add-on Credit Packs & Lifetime Deal Framework

- **Audit Credit Add-ons:** \$29 per pack of 50 additional deep domain audits.
- **Early-Adopter Lifetime Deal (LTD):** To build initial momentum and cash flow, a limited \$149 one-time LTD offer will be promoted via AppSumo and Product Hunt, capped at 200 tiers. This raises \$29,800 in non-dilutive capital to fund early API usage.

Strategic Recommendations

- Offer a 14-day free trial on the Pro tier to lower friction and drive self-serve conversions.
- Require annual prepayments on the Agency tier to secure cash flow upfront.

17. Multi-Channel Marketing and Launch Execution

The go-to-market plan combines developer community outreach with organic search acquisition.

Go-To-Market Channels:

— Product Hunt Launch	(Top-of-funnel tech awareness & LTD launch)
— Hacker News "Show HN"	(Developer-focused llms.txt validator showcase)
— Reddit Community Strategy	(Educational posts on r/SEO, r/SaaS, r/WebDev)
— Programmatic SEO Engine	(120,000+ domain scorecards indexed)
— Agency Affiliate Network	(30% recurring commissions for consultants)

Channel Execution Playbooks

- Product Hunt Launch:**
 - *Headline:* "AIQualityHQ — Measure Your Domain's AI Visibility & Validate llms.txt"
 - *Offer:* Free deep domain scorecard for all launch-day community members.
- Hacker News "Show HN":** * *Angle:* Focus on the technical implementation of the llms.txt specification and AI crawler behavior.
 - *Title:* "Show HN: Instant open-source validator for llms.txt and AI search crawlers".
- Community Outreach & Content Marketing:** * Publish diagnostic guides on Reddit (r/SEO, r/SaaS) explaining how to permit OAI-SearchBot while restricting GPTBot.
- Agency Affiliate Program:**
 - Offer a 30% recurring lifetime commission for SEO consultants reselling white-label scorecards to clients.

Strategic Recommendations

- Prepare all launch collateral (demo videos, screenshots, sample reports) at least two weeks before launch.
- Engage actively in launch-day comment threads to build community goodwill.

18. Programmatic Backlink Engine

To establish domain authority and earn natural backlinks, AIQualityHQ.com deploys four automated link-building mechanisms:

Viral Embeddable Badge Backlink Loop:

```
[ Domain Owner Audits Site ] —> [ Scores  
[span_373] (start_span) [span_373] (end_span) [span_378] (start_span) [span_  
378] (end_span)>85 Grade ] —> [ Embeds Badge on Site Footer ]
```

```
[span_423] (start_span) [span_423] (end_span) [span_427] (start_span) [span_  
427] (end_span)
```



[Visitors Click Badge] <— [Authority Transferred] <— [Link[span_26] (start_span) [span_26] (end_span) [span_46] (start_span) [span_46] (end_span) s to AIQualityHQ Scorecard]

1. **Embeddable "Verified AI-Ready" Dynamic Badges:** Domains scoring above 85/100 receive a dynamic SVG badge ("Verified AI-Ready by [span_143](start_span)[span_143](end_span)[span_154](start_span)[span_154](end_span)AIQualityHQ"). When embedded in a site's footer, it links back to their public scorecard (AIQualityHQ.com/score/domain.com), creating a continuous backlink loop.
2. **Global lms.txt Public Directory:** A public directory indexing verified lms.txt files across the web, earning citations from developer blogs and media outlets.
3. **Quarterly AI Search Visibility Reports:** Benchmark studies analyzing AI crawler access across the Fortune 500, distributed to digital PR networks.

Strategic Recommendations

- Ensure all embeddable badge code snippets use clean, non-intrusive HTML with descriptive alt text for SEO value.
- Automatically notify domain owners when their scorecard grade improves so they are encouraged to share their score publicly.

19. Financial Model and Projections

The business model delivers high gross margins enabled by serverless edge infrastructure and programmatic organic traffic growth.

Financial Metric	Year 1	Year 2	Year 3	Year 5
Monthly Organic Traffic	80,000 visits	350,000 visits	1,200,000 visits	3,500,000 visits
Free Registered Users	12,000 users	65,000 users	240,000 users	850,000 users
Paid Subscribers	350 customers	1,850 customers	5,400 customers	14,500 customers
Average Revenue Per User (ARPU)	\$58 / mo	\$64 / mo	\$72 / mo	\$85 / mo
Monthly Recurring Revenue (MRR)	\$20,300	\$118,400	\$388,800	\$1,232,500
Annual Recurring Revenue (ARR)	\$243,600	\$1,420,800	\$4,665,600	\$14,790,000
Infrastructure & API Costs	\$2,800 / mo	\$11,200 / mo	\$31,000 / mo	\$85,000 / mo
Gross Margin %	86.2%	90.5%	92.0%	93.1%
Net Founder Profit (EBITDA)	\$182,000	\$1,180,000	\$3,950,000	\$12,800,000

Strategic Recommendations

- Reinvest early profits into expanding the programmatic SEO index to accelerate organic traffic growth.
- Maintain a lean cost structure to ensure profitability even during traffic fluctuations.

20. Comprehensive Risk Assessment and Mitigation

Risk Category	Identified Risk Factor	Severity / Impact	Countermeasure & Mitigation Strategy
Technical / API	SERP/LLM providers block scrapers or increase API pricing significantly.	High	Maintain multi-provider failover routing (DataForSEO, Serper.dev, Direct Headless Scraping).
Market Standard	The llms.txt specification fails to achieve broad mainstream adoption.	Medium	Maintain product focus on overall AI Visibility, Entity Schema, and Multi-Engine Citation Tracking.
SEO Risk	Google penalizes large-scale programmatic SEO scorecard pages.	High	Ensure high unique value per page by injecting dynamic edge audit metrics, JSON-LD, and custom recommendations.
Legal / Privacy	AI platforms claim unauthorized prompt-monitoring or trademark issues.	Low	Utilize standard public API endpoints and process publicly accessible web data exclusively.
Competition	Competitors launch free instant scorecard utilities.	Medium	Build strong distribution moats through programmatic SEO index depth and open-source developer tooling.

Strategic Recommendations

- Audit programmatic SEO pages regularly to ensure they maintain high indexation rates and content quality standards.
- Stay updated on AI crawler policy changes to adjust diagnostic audit rules promptly.

21. 90-Day Execution Roadmap

Roadmap Overview:

└─ Weeks 1-3: Core Systems Engineering (DB schema, parser engines, edge setup)

|— Weeks 4-6: pSEO

Infrastructure[span_494] (start_span) [span_494] (end_span) & MVP Launch (pSEO pages, UI, Stripe billing)

|— Weeks 7-9: Public Launch & Growth (Product Hunt, Hacker News, Reddit)

|— Weeks 10-12: Monetization Scaling (Pro/Agency features, affiliate network)

- **Weeks 1–3: Core Systems & Engine Architecture**
 - Set up Supabase PostgreSQL schema, Next.js App Router, and Cloudflare Worker environments.
 - Build pure JavaScript AST parser for /llms.txt and /robots.txt AI user-agent parsing.
 - Integrate DataForSEO SERP API for automated citation querying.
- **Weeks 4–6: pSEO Infrastructure & Beta Deployment**
 - Implement dynamic page routes for /score/[domain] and /llms-txt/[domain].
 - Generate initial sitemap indexing 10,000 domain scorecards.
 - Connect Stripe Billing for Pro (\$49/mo) and Agency (\$199/mo) plans.
- **Weeks 7–9: Public Launch & Community Outreach**
 - Launch on Product Hunt; post technical showcase to Hacker News.
 - Publish educational guides across Reddit (r/SEO, r/SaaS).
- **Weeks 10–12: Feature Expansion & Path to \$10k MRR**
 - Deploy automated llms.txt file builder to increase lead conversions.
 - Ship Share of Voice competitor tracking for Agency subscribers.
 - Expand pSEO database to index 100,000 domains.

Strategic Recommendations

- Maintain a strict focus on launch deadlines, prioritizing core diagnostic features over optional enhancements.
- Collect user feedback actively during beta testing to refine audit scoring accuracy before public launch.

22. Final Founder Playbook

MVP Scope (Build First)

1. **Instant Domain Audit Engine:** Single-input search interface generating instant domain scorecards without requiring registration.
2. **llms.txt Syntax Validator:** Parser evaluating Markdown syntax against Jeremy Howard's specification.
3. **AI Crawler Permission Checker:** Diagnostic module checking robots.txt for OAI-SearchBot, Claude-SearchBot, and PerplexityBot access.
4. **Programmatic SEO Engine:** Deploy top 10,000 domain scorecards at /score/[domain].

Non-Essential Features (De-prioritize Initially)

- **Custom Server Log Monitoring:** Delay enterprise log ingestion features until reaching

\$20k MRR.

- **Automated Content Rewriting:** Focus on diagnostic scoring accuracy before building complex content generation suites.

Monetization Path (\$0 to \$100k MRR)

Monetization Milestones:

\$0 [span_27] (start_span) [span_27] (end_span) [span_47] (start_span) [span_47] (end_span) MRR —————> \$10k MRR —————> \$50k MRR

—————> \$100k MRR

Product Hunt & 100k pSEO Pages 1,000 Pro Users
Agency Scaling

Hacker News Launch + Lifetime Deals + 50 Agencies
[span_270] (start_span) [span_270] (end_span) [span_277] (start_span) [span_277] (end_span) + Enterprise Contracts

1. **Phase 1 (\$0 to \$10k MRR):** Execute Product Hunt launch, collect capital via early Lifetime Deals (\$149), and index 50,000 pSEO domain scorecards.
2. **Phase 2 (\$10k to \$50k MRR):** Expand pSEO index to 250,000 pages, leverage embeddable badges to build domain authority, and convert agencies to \$199/mo plans.
3. **Phase 3 (\$50k to \$100k MRR):** Introduce enterprise log monitoring, expand integration partnerships (Vercel, Mintlify, Cloudflare), and solidify AIQualityHQ.com as the industry standard for AI search visibility.

Works cited

1. Best AI Visibility Tools 2026: Profound vs Peec vs Otterly vs the Rest | Surmado Blog, <https://www.surmado.com/blog/best-ai-visibility-tools-2026> 2. What Is LLMS.txt And Ways To Create An LLMS.txt File - IDCM, <https://www.myidcm.com/blog/llms-txt> 3. AI Crawlers Explained: GPTBot, ClaudeBot, PerplexityBot and How to Let Them In (2026), <https://www.anagram.ai/blog/ai-crawlers-explained-gptbot-claudebot-perplexitybot-and-how-to-let-them-in-2026> 4. Profound vs Peec AI (2026): Which AI Visibility Tool Should You Pick? | Duple, <https://duple.com/learn/profound-vs-peec-ai> 5. 9 Top Profound AI Alternatives in 2026 for AI Visibility and Brand Monitoring - Reddit, https://www.reddit.com/r/geotoolsreview/comments/1v57q5h/9_top_profound_ai_alternatives_in_2026_for_ai/ 6. Peec AI - AI Search Analytics for Marketing Teams, <https://peec.ai/> 7. The 8 best AI visibility tools in 2026 - Zapier, <https://zapier.com/blog/best-ai-visibility-tool/> 8. Llms.txt: The Complete Guide: What It Is, How to Create One, and Why It Matters for AI Search - Reaudit, <https://reaudit.io/blog/llms-txt-guide-complete-tutorial> 9. What is Llms.txt? Breaking down the skepticism - Mintlify, <https://www.mintlify.com/blog/what-is-llms-txt> 10. //lms.txt—a proposal to provide information to help LLMs use websites – Answer.AI, <https://www.answer.ai/posts/2024-09-03-llmstxt.html> 11. Llms.txt: The //lms.txt file, <https://llmstxt.org/> 12. AI Crawler Directory - DataFast, <https://datafa.st/crawlers> 13. AI Crawler User-Agent List | GPTBot, ClaudeBot, PerplexityBot & More - ClickFrom.AI, <https://www.clickfrom.ai/tools/ai-crawler-user-agent-list> 14. How to scrape Google AI Overviews with SERP API - DataForSEO, <https://dataforseo.com/help-center/how-to-scrape-google-ai-overviews-with-serp-api> 15. AI Agents for Search Marketers - BrightEdge,

<https://www.brightedge.com/resources/guide-for-ai-agents> 16. Overview of OpenAI Crawlers, <https://developers.openai.com/api/docs/bots> 17. List of Top AI Search Crawlers + User Agents (Winter 2025) | Momentic, <https://momenticmarketing.com/blog/ai-search-crawlers-bots>